


ITE / TTE vs. system size, all archived solvers, each at its own best-available operating point
 $h=0.5$ matched across all series; lr/reg/n_samples/iter budget still differ -- see script header for each series' config, N range, n
 (hollow markers at the same x are nudged apart vertically to stay distinguishable; the offset is cosmetic, not data)

ITE — iterations to 0.01 energy error/spin (median, IQR; hollow = censored)

(A) Iterations-to-epsilon vs. N



(B) Time-to-epsilon vs. N (D-Wave QPU shown only where timing was recorded)

